

Lab Instructions - session 5

Reading from camera device, edge detection

Part 1. Template Matching

Template matching is a technique used to find small parts of an image that match a template. This method uses correlation to compare the template against overlapping regions of the scene. `cv2.matchTemplate()` slides the template across the image and computes similarity scores.

File: `match.py`

```
import cv2
from matplotlib import pyplot as plt

# Load the grayscale scene image
I = cv2.imread("scene.png", cv2.IMREAD_GRAYSCALE)

# Load the grayscale template image
template = cv2.imread("template.png", cv2.IMREAD_GRAYSCALE)

# Compute correlation between template and image
res = cv2.matchTemplate(I, template, cv2.TM_CCOEFF)

# Find the location of the best match
min_val, max_val, min_loc, max_loc = cv2.minMaxLoc(res)

# Get template width and height
w, h = template.shape[::-1]

# Define top-left and bottom-right coordinates of the matched region
top_left = max_loc
bottom_right = (top_left[0] + w, top_left[1] + h)

# Draw a rectangle around the detected template location
cv2.rectangle(I, top_left, bottom_right, 255, 2)

# Display the correlation map and detected template location
plt.subplot(121), plt.imshow(res, cmap='gray')
plt.title('Matching Result'), plt.xticks([]), plt.yticks([])
plt.subplot(122), plt.imshow(I, cmap='gray')
plt.title('Detected Point'), plt.xticks([]), plt.yticks([])

plt.show()
```

- What does `cv2.matchTemplate()` do, and how does it use correlation for template matching?
- What happens if the template appears multiple times in the scene?

Part 2. Reading from a webcam

The following code reads a video stream from a webcam and displays it.

File: `read_webcam.py`

```
import numpy as np
import cv2

cam_id = 0 # camera id

# for default webcam, cam_id is usually 0
# try out other numbers (1,2,...) if this does not work
# If you are on linux, you can use the command
# "ls /dev/video*" to see available webcams.
cap = cv2.VideoCapture(cam_id)

while True:
    ret, I = cap.read();

    cv2.imshow("my stream", I);

    # press "q" to quit
    if cv2.waitKey(1) & 0xFF == ord('q'):
        break

cap.release()
cv2.destroyAllWindows()
```

Part 3. Edge detection

Computing the Gradients

We compute Sobel gradients both by applying a Sobel filter and by using the OpenCV function “Sobel”.

File: **sobel1.py**

```
import numpy as np
import cv2
from matplotlib import pyplot as plt

I = cv2.imread("agha-bozorg.jpg", cv2.IMREAD_GRAYSCALE)

# Compute the gradient in x direction using the sobel filter

# Method 1: using filter2D *****
Dx = np.array([[-1, 0, 1],
               [-2, 0, 2],
               [-1, 0, 1]]) # Sobel filter

Ix = cv2.filter2D(I, -1, Dx);
print(I.dtype)
print(Ix.dtype)

Ix = cv2.filter2D(I, cv2.CV_16S, Dx) # cv2.CV_16S: 16 bit signed integer
print(Ix.dtype)
input('press ENTER to continue... ')

# Method 2: using sobel function *****
Ix2 = cv2.Sobel(I,cv2.CV_16S,1,0)
print(np.abs(Ix - Ix2).max())
input('press ENTER to continue... ')

# Plot the gradient image
f, axes = plt.subplots(2, 2)

axes[0,0].imshow(I,cmap = 'gray')
axes[0,0].set_title("Original Image")

axes[0,1].imshow(Ix,cmap = 'gray')
axes[0,1].set_title("Ix (cv2.filter2D)")

axes[1,0].imshow(Ix2,cmap = 'gray')
axes[1,0].set_title("Ix2 (cv2.Sobel)")

axes[1,1].imshow(np.abs(Ix),cmap = 'gray')
axes[1,1].set_title("abs(Ix)")

# Notice that imshow in matplotlib considers the minimums value of I
# as black and the maximum value as white (this is different from
# the behavior in cv2.imshow
plt.show()
```

- Why have we used `cv2.CV_16S` (16 bit signed integer) format for the output of `filter2D` and `Sobel` functions, instead of `-1` (which gives the same numeric type as the input image, i.e. `CV_8U` or unsigned 8-bit integer)?
- What is the value of `np.abs(Ix - Ix2).max()`? Are `Ix` and `Ix2` different?

- Why are most of the pixels in images of `axes[0,1]` and `axes[1,0]` gray? What do the black and white pixels show in these images? (notice that matplotlib automatically sets the minimum value of an image to black and the maximum value to white)

Task 1:

Modify `sobel1.py` to include Prewitt and Roberts filters alongside the existing Sobel filter and apply all three filters to `edge.png`. Then display and compare the results.

- How do the detected edges differ between Sobel, Prewitt, and Roberts filters?
- How does the size of the convolution kernel affect edge detection?
- Which filter introduces more noise in the output? Explain your observation.

Towards edge detection!

The strength of the edge at each certain pixel can be represented by the magnitude of the gradient vector, that is

$$\sqrt{(\partial I / \partial x)^2 + (\partial I / \partial y)^2}$$

The next code plots the gradient in x and y directions, plus the magnitude of the gradient.

File: `sobel2.py`

```
import numpy as np
import cv2
from matplotlib import pyplot as plt

I = cv2.imread("agha-bozorg.jpg", cv2.IMREAD_GRAYSCALE)

# Sobel gradient in the x direction
Ix = cv2.Sobel(I, cv2.CV_64F, 1, 0)
print(Ix.dtype)

# Sobel gradient in the y direction
Iy = cv2.Sobel(I, cv2.CV_64F, 0, 1)
print(Iy.dtype)

# Magnitude of gradient
E = np.sqrt(Ix*Ix + Iy*Iy)
```

```
# Plot the gradient image
f, axes = plt.subplots(2, 2)

axes[0,0].imshow(I,cmap = 'gray')
axes[0,0].set_title("Original Image")

axes[0,1].imshow(abs(Ix),cmap = 'gray')
axes[0,1].set_title("abs(Ix)")

axes[1,0].imshow(abs(Iy),cmap = 'gray')
axes[1,0].set_title("abs(Iy)")

axes[1,1].imshow(E,cmap = 'gray')
axes[1,1].set_title("Magnitude of Gradient")
plt.show()
```

Part 4. Recall images as functions

Derivatives and edges

In this task, you will explore how images can be represented as 3D functions, where pixel intensities define a surface. You will generate surface plots for:

1. The original grayscale image
2. The Gaussian-blurred image
3. The Second derivative of Gaussian (Laplacian of Gaussian, LoG)

The function **zero_crossing** finds zero crossings in an image for LoG edge detection. You do not need to know how the functions **std_filter** and **zero_crossing** work.

File: **LoG.py**

```
import numpy as np
import cv2
from matplotlib import pyplot as plt

def std_filter(I, ksize):
    F = np.ones((ksize,ksize), dtype=np.float) / (ksize*ksize)
    MI = cv2.filter2D(I,-1,F) # apply mean filter on I
    I2 = I * I; # I squared
    MI2 = cv2.filter2D(I2,-1,F) # apply mean filter on I2
    return np.sqrt(abs(MI2 - MI * MI))

def zero_crossing(I):
    """Finds locations at which zero-crossing occurs, used for
```

```
Laplacian edge detector"""

Ishrxx = I.copy();
Ishrxx[:,1:] = Ishrx[:, :-1]
Ishdy = I.copy();
Ishdy[1:,:] = Ishdy[:-1,:]
ZC = (I==0) | (I * Ishrx < 0) | (I * Ishdy < 0); # zero crossing
locations
SI = std_filter(I, 3) / I.max()
Mask = ZC & (SI > .1)
E = Mask.astype(np.uint8) * 255 # the edges
return E

# Load the image in grayscale
image = cv2.imread("cameraman.png", cv2.IMREAD_GRAYSCALE)

blurred_image = cv2.GaussianBlur(image, (7, 7), 0)
El = cv2.Laplacian(blurred_image, cv2.CV_64F, ksize=5)
log_image = zero_crossing(El)

# Get image dimensions
rows, cols = image.shape
X, Y = np.meshgrid(np.arange(cols), np.arange(rows))

# Create figure with three subplots
fig = plt.figure(figsize=(18, 6))

# Plot Original Image Surface
ax1 = fig.add_subplot(131, projection='3d')
ax1.plot_surface(X, Y, image, cmap='gray', edgecolor='none')
ax1.set_title("Original Image")
ax1.set_xlabel("X axis (Width)")
ax1.set_ylabel("Y axis (Height)")
ax1.set_zlabel("Pixel Intensity")

# Plot Gaussian Blurred Image Surface
ax2 = fig.add_subplot(132, projection='3d')
ax2.plot_surface(X, Y, blurred_image, cmap='gray', edgecolor='none')
ax2.set_title("Gaussian Blurred Image")
ax2.set_xlabel("X axis (Width)")
ax2.set_ylabel("Y axis (Height)")
ax2.set_zlabel("Pixel Intensity")

# Plot Second Derivative (Laplacian of Gaussian)
ax3 = fig.add_subplot(133, projection='3d')
ax3.plot_surface(X, Y, log_image, cmap='gray', edgecolor='none')
ax3.set_title("Second Derivative (LoG)")
ax3.set_xlabel("X axis (Width)")
ax3.set_ylabel("Y axis (Height)")
ax3.set_zlabel("Pixel Intensity")

# Show the plots
plt.show()
```

- How does the Gaussian blur affect the 3D surface plot? How does increasing or decreasing the Gaussian kernel's size change the surface plot's smoothness?
- What differences do you observe in the second derivative (LoG) plot compared to the original image?
- How can you identify edges in the 3D surface plot of the original image and the Laplacian of Gaussian plot? What characteristics indicate edge transitions?

Edge detection (Sobel, Laplacian, Canny):

Now, we move on to edge detection. We use three different methods:

1. Using Sobel gradients + thresholding
2. Laplacian of Gaussian (LoG) + detection zero-crossings
3. Canny Edge detection

File: **edge.py**

```
import numpy as np
import cv2
from matplotlib import pyplot as plt

def std_filter(I, ksize):
    F = np.ones((ksize,ksize), dtype=np.float) / (ksize*ksize);

    MI = cv2.filter2D(I,-1,F) # apply mean filter on I

    I2 = I * I; # I squared
    MI2 = cv2.filter2D(I2,-1,F) # apply mean filter on I2

    return np.sqrt(abs(MI2 - MI * MI))

def zero_crossing(I):
    """Finds locations at which zero-crossing occurs, used for
    Laplacian edge detector"""

    Ishrx = I.copy();
    Ishrx[:,1:] = Ishrx[:,:,:-1]

    Ishdy = I.copy();
    Ishdy[1:,:] = Ishdy[:-1,:]

    ZC = (I==0) | (I * Ishrx < 0) | (I * Ishdy < 0); # zero crossing
    locations

    SI = std_filter(I, 3) / I.max()

    Mask = ZC & (SI > .1)

    E = Mask.astype(np.uint8) * 255 # the edges
```

```
return E

I = cv2.imread("agha-bozorg.jpg", cv2.IMREAD_GRAYSCALE)

# set the sigma for Gaussian Blurring
sigma = 7

# Sobel magnitude of gradient
thresh = 90 # threshold
Ib = cv2.GaussianBlur(I, (sigma,sigma), 0); # blur the image
Ix = cv2.Sobel(Ib,cv2.CV_64F,1,0)
Iy = cv2.Sobel(Ib,cv2.CV_64F,0,1)
Es = np.sqrt(Ix*Ix + Iy*Iy)
Es = np.uint8(Es > thresh)*255 # threshold the gradients

# Laplacian of Gaussian
# Here, we first apply a Gaussian filter and then apply
# the Laplacian operator (instead of applying the LoG filter)
Ib = cv2.GaussianBlur(I, (sigma,sigma), 0);
El = cv2.Laplacian(Ib,cv2.CV_64F,ksize=5)
El = zero_crossing(El);

# Canny Edge detector
lth = 50    # low threshold
hth = 120   # high threshold
Ib = cv2.GaussianBlur(I, (sigma,sigma), 0); # blur the image
Ec = cv2.Canny(Ib,lth, hth)
f, axes = plt.subplots(2, 2)

axes[0,0].imshow(I,cmap = 'gray')
axes[0,0].set_title("Original Image")

axes[0,1].imshow(Es,cmap = 'gray')
axes[0,1].set_title("Sobel")

axes[1,0].imshow(El,cmap = 'gray')
axes[1,0].set_title("Laplacian")

axes[1,1].imshow(Ec,cmap = 'gray')
axes[1,1].set_title("Canny")

# Notice that imshow in matplotlib considers the minimums value of I
# as black and the maximum value as white (this is different from
# the behavior in cv2.imshow)
plt.show()
```

- Compare the Canny edge detector to Sobel+thresholding. Can you see the effect of non-maximum suppression?
- Notice that in all cases we first smooth the image using a Gaussian filter. What is the purpose of smoothing the image? Change the smoothing parameter **sigma**



(from **sigma=7** to sigma=5, 3, 1 or 9) and see what happens by increasing or decreasing this parameter.

- Change the low and high thresholds of the Canny edge detector and observe the results.

Task 2:

You need to read a video stream from your webcam and apply different gradients or edge detection operations to the stream. You have to do this by completing the file **lab5_task1.py**. Your program must have the following functionalities:

- **press the 'o' key:** show the original webcam frame (already done)
- **press the 'x' key:** show the Sobel gradient in the x direction (already done)
- **press the 'y' key:** show the Sobel gradient in the y direction
- **press the 'm' key:** show the Sobel magnitude of the gradient
- **press the 's' key:** show the result of Sobel + thresholding edge detection
- **press the 'l' key:** apply Laplacian of Gaussian (LoG) edge detector
- **press the 'c' key:** apply Canny edge detector
- **press the '+' key:** increase smoothing parameter sigma (already done)
- **press the '-' key:** decrease smoothing parameter sigma (already done)
- **press the 'q' key:** quit the program (already done)

Notice that after reading each image frame we apply a Gaussian filter to blur it. All gradient/edge detection operations must be done on the blurred image **Ib**.

File: **lab5_task2.py**

```
import numpy as np
import cv2

cam_id = 0 # camera id

# for default webcam, cam_id is usually 0
# try out other numbers (1,2,...) if this does not work
#
cap = cv2.VideoCapture(cam_id)

mode = 'o' # show the original image at the beginning

sigma = 5

while True:
    ret, I = cap.read();
    #I = cv2.imread("agha-bozorg.jpg") # can use this for testing
```

```
I = cv2.cvtColor(I, cv2.COLOR_BGR2GRAY) # convert to grayscale
Ib = cv2.GaussianBlur(I, (sigma,sigma), 0); # blur the image

if mode == 'o':
    # J = the original image
    J = I
elif mode == 'x':
    # J = Sobel gradient in x direction
    J = np.abs(cv2.Sobel(Ib,cv2.CV_64F,1,0));

elif mode == 'y':
    # J = Sobel gradient in y direction
    pass
elif mode == 'm':
    # J = magnitude of Sobel gradient
    pass
elif mode == 's':
    # J = Sobel + thresholding edge detection
    pass
elif mode == 'l':
    # J = Laplacian edges
    pass
elif mode == 'c':
    # J = Canny edges
    pass

# we set the image type to float and the maximum value to 1
# (for a better illustration) notice that imshow in opencv does not
# automatically map the min and max values to black and white.
J = J.astype(np.float) / J.max();
cv2.imshow("my stream", J);

key = chr(cv2.waitKey(1) & 0xFF)

if key in ['o', 'x', 'y', 'm', 's', 'c', 'l']:
    mode = key
if key == '-' and sigma > 1:
    sigma -= 2
    print("sigma = %d"%sigma)
if key in ['+', '=']:
    sigma += 2
    print("sigma = %d"%sigma)
elif key == 'q':
    break

cap.release()
cv2.destroyAllWindows()
```

- Wave your hand quickly. What happens to the edges? Why?
- Press the 's' and 'c' keys to compare **Canny** with **Sobel** edge detection. Explain the effect of non-maximum suppression to the TA.

- In each case, increase and decrease the **sigma** by pressing + and - keys and see what happens. Explain the reason to the TA.
- Ideally, you can also use [`cv2.putText\(\)`](#) function to label your images so that you can compare them easily.

Task 3: Color Edge Detection in RGB Space

Traditional edge detection methods often rely on grayscale images, losing valuable color information. This task follows the approach from the article, which detects edges directly in the RGB space by considering the maximum directional differences across color channels.

The key steps of the algorithm are:

- Noise Reduction: The image is preprocessed using an adaptive median filter to reduce noise while preserving edges.
- Gradient Computation in RGB Space: Instead of converting the image to grayscale, Sobel filters are applied separately to the R, G, and B channels to compute gradients.
- Maximum Gradient Selection: At each pixel, the maximum gradient magnitude across the three channels is selected to form an edge map.
- Thresholding: A threshold is applied to retain strong edges while suppressing weak ones.

By implementing this approach, we retain color-based edge information, which can be useful for detecting object boundaries that are indistinguishable in grayscale images.

Now, complete the file `lab5_task3.py` to implement the algorithm.

File: `lab5_task3.py`

```
from matplotlib import pyplot as plt
import numpy as np
import cv2

def adaptive_median_filter(img, kernel_size=3):
    """Applies an adaptive median filter to reduce noise while preserving edges."""
    return cv2.medianBlur(img, kernel_size)
```

```
def max_directional_difference(img):
    """Computes the maximum directional difference in the RGB space."""
    diff = np.zeros(img.shape[:2], dtype=np.float32)
    for i in range(3): # Iterate over R, G, B channels
        #Compute Sobel in the x direction
        sobelx =
        #Compute Sobel in the y direction
        sobely =
        diff = np.maximum(diff, np.sqrt(sobelx**2 + sobely**2))
    return diff

def threshold_edge_detection(diff, threshold_ratio=0.3):
    """Applies a threshold to the detected edges."""
    threshold = threshold_ratio * np.max(diff)
    return (diff > threshold).astype(np.uint8) * 255

# Load image
image = cv2.imread("pepper.jpg")
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

# Apply adaptive median filter
filtered_img =

# Compute max directional difference
edge_map =

# Apply thresholding
final_edges =

# Display results
fig, ax = plt.subplots(1, 3, figsize=(15, 5))
ax[0].imshow(image)
ax[0].set_title("Original Image")
ax[0].axis("off")

ax[1].imshow(edge_map, cmap="gray")
ax[1].set_title("Max Directional Difference")
ax[1].axis("off")

ax[2].imshow(final_edges, cmap="gray")
ax[2].set_title("Final Edge Map")
ax[2].axis("off")

plt.show()
```

- How does the adaptive median filter affect the final edge detection result?
- What advantages does this method have over traditional grayscale-based edge detection?
- How does the choice of the threshold ratio impact the edge map?



References

- https://docs.opencv.org/3.0-beta/doc/py_tutorials/py_imgproc/py_gradients/py_gradients.html#gradients
- https://docs.opencv.org/3.0-beta/doc/py_tutorials/py_imgproc/py_canny/py_canny.html#canny